


System Programming Lab Session #3

Kernel Lab

2019/10/01

sysprog@csap.snu.ac.kr



Computer Systems and Platforms Laboratory
School of Computer Science and Engineering
Seoul National University

Goals

- Understand the art of the Linux Kernel
 - Learn how to program a kernel module in Linux
 - Understand the hierarchical structure of page tables and address translation

1. Introduction

What is a Kernel Module

- A module is a piece of code that can be loaded and unloaded into the kernel upon demand.
- Can use privileged instructions without system calls, because a kernel module is loaded and executed within the kernel.
- Module load / unload commands in Linux
 - Load: `root # insmod < module_name.ko >`
 - Unload: `root # rmmod <module_name>`
 - Module list `devel $ lsmod`

Linux Kernel Module Programming

- There are some conventions when programming a kernel module in Linux
- Initialize and Exit module
 - `module_init()`: Called when the module is inserted
 - `module_exit()`: Called when the module is removed

```
#include <linux/module.h>

MODULE_LICENSE("GPL");

static int __init init_my_module(void)
{
    // Running when this module is inserted to Kernel
}

static void __exit exit_my_module(void)
{
    // Running when this module is removed from Kernel
}

module_init(init_my_module);
module_exit(exit_my_module);
```

Linux Kernel Module Programming

- File operations
 - File operations are used to communicate with files in Device Driver and Debug File System
 - Function pointer structure for the file

```
struct file_operations Fops = {  
    .read = file_read,  
    .write = file_write,  
    .open = file_open,  
    .release = file_close,  
};
```

Linux Kernel Module Programming

- **Debug File System (debugfs)** is a special file system available in the Linux kernel. Debugfs is a simple-to-use RAM-based file system specially designed for debugging purposes. It exists as a simple way for the kernel developers to make information quickly and easily available to user space.
- **Debugfs** has no rules at all. Developers can put any type of information that they want.
- **Debugfs** also supports simple user-to-kernel interfaces in Linux kernel modules. Developers can access Linux kernel information easily using debugfs.

Debugfs APIs

- Code using debugfs must include `<linux/debugfs.h>`
- Debugfs APIs
 - **struct dentry*** debugfs_create_dir(**const char** *name, **struct dentry** *parent)
 - **struct dentry*** debugfs_create_file(**const char** *name, **struct dentry** *parent, **void** *data, **const struct file_operations** *fops)
 - **struct dentry*** debugfs_remove_recursive(**struct dentry** *dentry)

Debugfs APIs

- Make a source file *dbfs.c*
- Make a build script *Makefile*
- Build and Insert your module

devel \$ make

devel \$ sudo insmod dbfs.ko

- Check files in the debugfs

root # cd /sys/kernel/debug/dir

root # ls

```
#include <linux/module.h>
#include <linux/debugfs.h>

MODULE_LICENSE("GPL");

static ssize_t write_fop(...)
{
    ...
}

static struct file_operations dbfs_ops = {
    .write = write_fop,
};

static int __init init_dbfs_module(void)
{
    ...
}

static void __exit exit_dbfs_module(void)
{
    ...
}

module_init(init_dbfs_module);
module_exit(exit_dbfs_module);
```

HowTo: Makefile

- GNU *make* is one many build systems that keep track of how to build your program from the sources
- Makefile is the file that contains the instructions for *make*.

example) **devel** \$ **make**

```
devel@gentoo ~ $ make
make -C /lib/modules/4.9.6-gentoo-r1/build M=/home/devel modules;
make[1]: Entering directory '/usr/src/linux-4.9.6-gentoo-r1'
  CC [M]  /home/devel/dbfs_ptree.o
  Building modules, stage 2.
  MODPOST 1 modules
  CC      /home/devel/dbfs_ptree.mod.o
  LD [M]  /home/devel/dbfs_ptree.ko
make[1]: Leaving directory '/usr/src/linux-4.9.6-gentoo-r1'
sudo insmod dbfs_ptree.ko
```

HowTo: Makefile

- Simple Example

```
all : module_trigger benchmark driver

module_trigger :
    gcc -o module_trigger -I./module module_trigger.c

benchmark :
    cd benchmarks; make; mv *_banchmark ../

driver :
    cd module; make

clean :
    rm module_trigger *_benchmark; cd module; make clean
```

- all : ...
 - Defines targets. Each targets must be defined in the followings.
 - When you execute *make* without a target, the first defined target is built.
- clean :
 - Defines removal of all built object files.
 - You can call it as 'make clean'

HowTo: Makefile for Kernel Modules

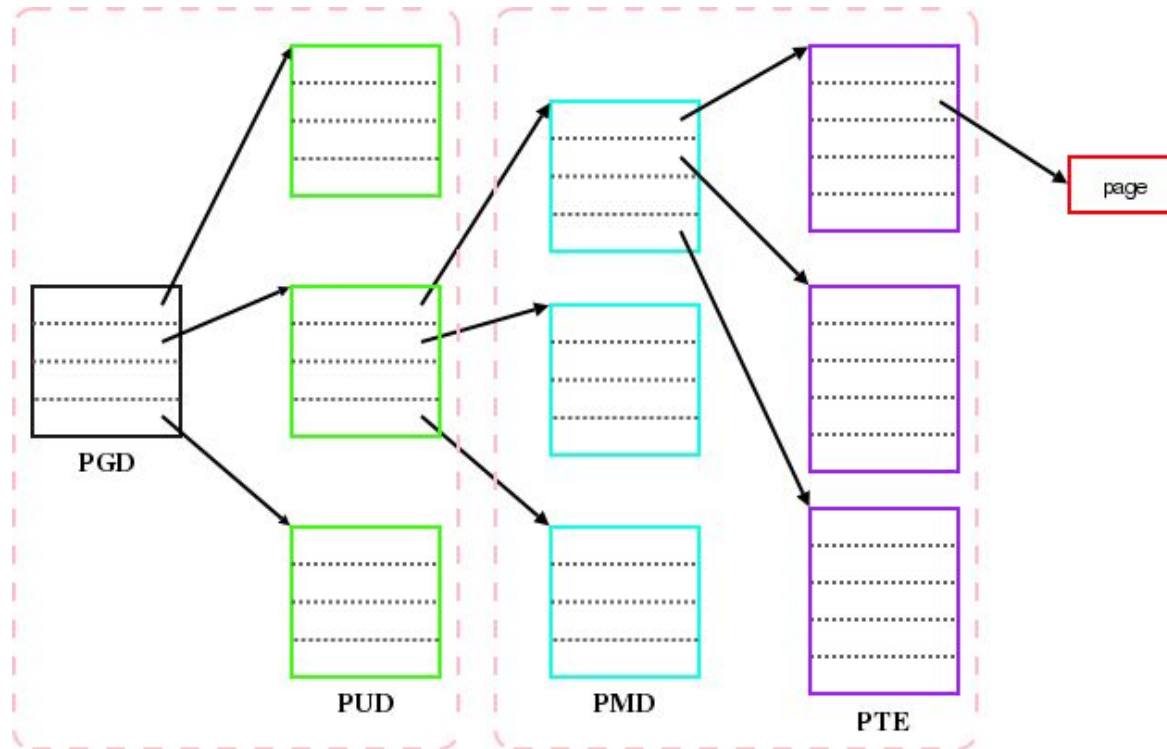
- Module Compile Example

```
obj-m := my_mod.o  
  
all :  
    $(MAKE) -C /lib/modules/$(shell uname -r)/build M=$(PWD) modules;  
    sudo insmod my_mod.ko  
  
clean :  
    $(MAKE) -C /lib/modules/$(shell uname -r)/build M=$(PWD) clean;  
    sudo rmmod my_mod.ko
```

- Kernel module is not compiled with general 'gcc'. It needs to be built using kernel-specific compile tools.
- The example shows a simple makefile for a kernel module build.
- You should change the module name in the red box.

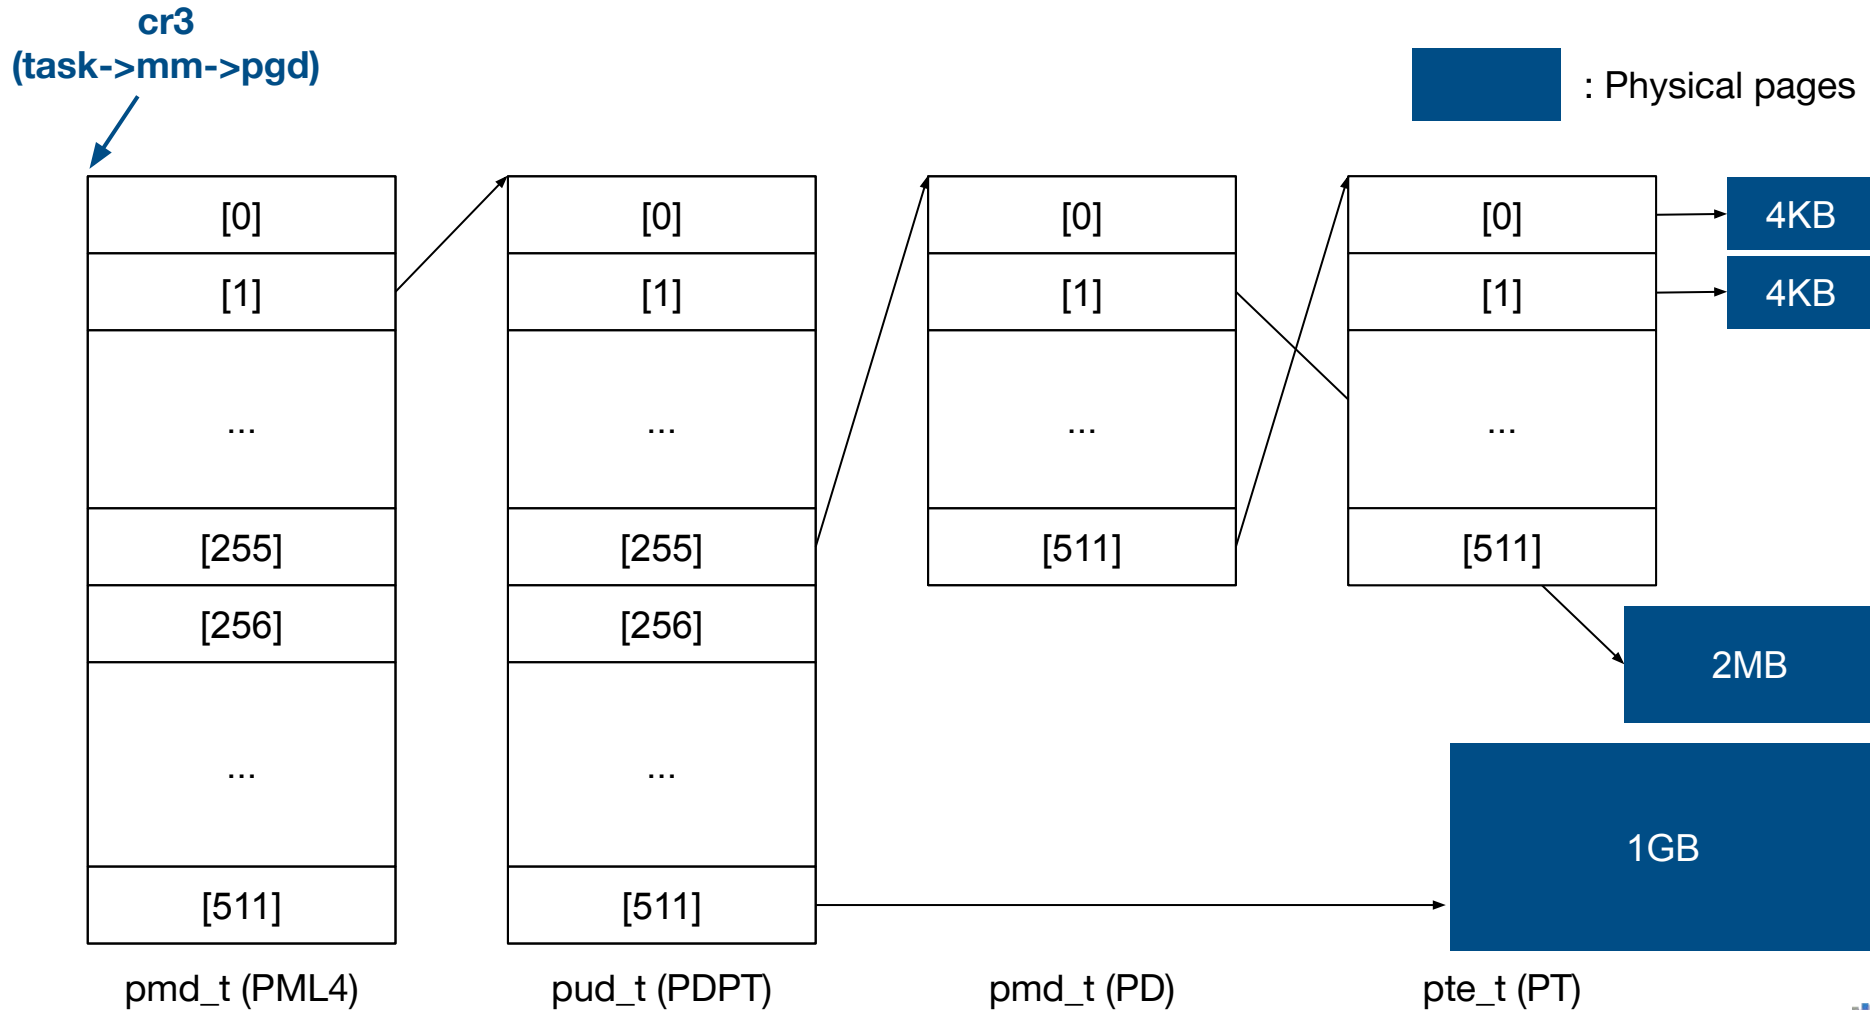
2. Background

Paging on a 64-bit machine




Paging on a 64-bit machine

- Typical page size is 4 KB, however there are some pages with 1 GB/2 MB size
- Page tables resident in 4 KB pages (4 kB / 8 B = 512 entries per table)



Paging on a 64-bit machine

- 512 entries = 2^9 entries

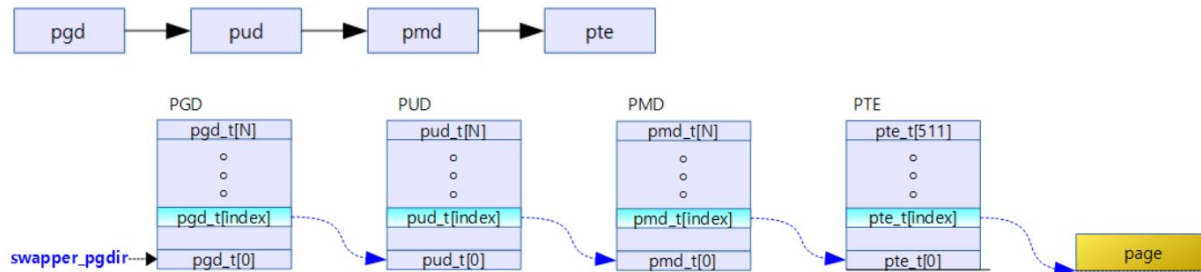


- Linear address space: 2^{48} Bytes = 256 TB
- Although virtual addresses are 64 bits wide, current implementations do not allow the entire virtual address space of 2^{64} bytes (16 EB) to be used.
- 0x00000000 00000000 ~ 0x00007FFF FFFFFFFF : Lower half (user space, 128 TB)
 - 0000000000000000 00000000 00000000 00000000 00000000 000000000000
 - 0000000000000000 01111111 11111111 11111111 11111111 111111111111
- 0xFFFF8000 00000000 ~ 0xFFFFFFFF FFFFFFFF : Higher half (kernel space, 128 TB)
 - 1111111111111111 10000000 00000000 00000000 00000000 000000000000
 - 1111111111111111 11111111 11111111 11111111 11111111 111111111111

3. Part A

Part A : Find Physical Address

- Your task: print the physical address for a given virtual address in the current process
- Use `getpid()` to get the PID of the current process
- For verification purposes, we use `mmap()` to fix a pointer to a given virtual address
- Include page walk procedure in your program



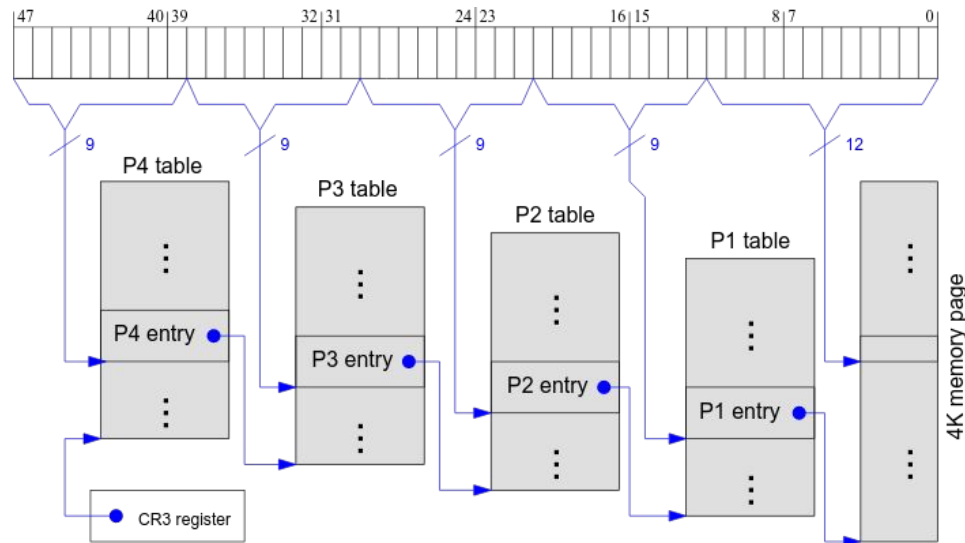
- **Sample Output**

```
devel@gentoo ~/kernellab/paddr $ sudo ./va2pa
[TEST CASE]      PASS
```

Part A : Find Physical Address

- Hints

- You can find the page walk APIs in `/usr/src/linux/arch/x86/include/asm/pgtable.h`
- Basic conversion from virtual address to physical address
- Access to the next level page table needs to be done in virtual address
 - ✓ Apply this conversion scheme to hierarchical page table.



Part A : Find Physical Address

- *make* command include *insmod*.
 - Include *insmod* under the *all* section in your *Makefile*.
- Do not change the debugfs directory and file name.
 - Directory name: *paddr*
 - Physical address output file name: *output*
- Do not modify the file *va2pa.c*
 - We will test your program using uniform *va2pa.c*.
 - Just use *va2pa.c* for testing your implementation.

4. Part B

Part B : Calculate the Resident Set Size of a Process

- **Resident set size (RSS)** is the portion of memory occupied by a process on the memory
- Print the *resident set size* of *test* program
- First, execute the program *test* and get its pid.
- Implement your kernel module to walk through the page tables of *test* program and print result
- **Sample Output**

```
devel@gentoo ~/kernellab/ptrav $ ./test 1000000
```

```
My PID: 5732
```

```
Allocated: 8000000 Bytes
```

```
devel@gentoo ~/kernellab/ptrav $ sudo ./rss 5732
```

```
1G pages: 0, 2M pages: 3, 4K pages: 606
```

```
Resident Set Size: 8568 kB
```

Part B : Calculate the Resident Set Size of a Process

- **Hints**

- You can find the relevant APIs in `/usr/src/linux/arch/x86/include/asm/pgtable.h`
- `task->mm->pgd` contains the address of first level page table (pgd or PML4)
- Page size of a page table is 4KB, how many entries in 64-bit?
- Virtual address space (user): `0x0000000000000000~0x00007FFFFFFFFFFFFFFF` (lower 128TB)
- Refer to the Intel Software Developer's Manual
 - ✓ **Examine which bits are used for which purposes (present, page size, etc.)**

Part B : Calculate the Resident Set Size of a Process

6	6	6	6	5	5	5	5	5	5	5	5	M ¹	M-1			3	3	3	2	2	2	2	2	2	2	2	1	1	1	1	1	1	1	1	1	0	
3	2	1	0	9	8	7	6	5	4	3	2	1			2	1	0	9	8	7	6	5	4	3	2	1	0		P	P	C	W	I	gn.	CR3		
Reserved ²													Address of PML4 table													Ignored			P C W D T	I gn.	CR3						
X D 3	Ignored											Rsvd.	Address of page-directory-pointer table													Ign.	R s v d	I g n	P A C E	P C W D T	U / S / W	R / W	1	PML4E: present			
Ignored																											Q		PML4E: not present								
X D 3	Prot. Key ⁴	Ignored				Rsvd.	Address of 1GB page frame				Reserved				P A T	Ign.	G	1	D	A	P C W D T	P C W D T	U / S / W	R / W	1	PDPTE: 1GB page											
X D 3	Ignored					Rsvd.	Address of page directory													Ign.	Q	I g n	A	P C W D T	P C W D T	U / S / W	R / W	1	PDPTE: page directory								
Ignored																											Q		PDPTE: not present								
X D 3	Prot. Key ⁴	Ignored				Rsvd.	Address of 2MB page frame				Reserved				P A T	Ign.	G	1	D	A	P C W D T	P C W D T	U / S / W	R / W	1	PDE: 2MB page											
X D 3	Ignored					Rsvd.	Address of page table													Ign.	Q	I g n	A	P C W D T	P C W D T	U / S / W	R / W	1	PDE: page table								
Ignored																											Q		PDE: not present								
X D 3	Prot. Key ⁴	Ignored				Rsvd.	Address of 4KB page frame													Ign.	G	P A T	D	A	P C W D T	P C W D T	U / S / W	R / W	1	PTE: 4KB page							

Part B : Calculate the Resident Set Size of a Process

- *make* command include *insmod*.
 - Include *insmod* under the *all* section in your *Makefile*.
- Do not change the debugfs directory and file name.
 - Directory name: *ptrav*
 - Physical address output file name: *output*
- Never modify the file *rss.c* or *test.c*
 - We will test your program using them.
 - Just use *rss.c* & *test.c* for testing your implementation.

5. Report

Report

- Your report should not be longer than 6 pages (excluding the cover page).
- Avoid copy-pasting screenshots of your code. We have your code. What we ask for here is your thought process applied to solve the lab. (More of a general remark since you are not submitted code but the idea remains)
- You can also attach some diagrams to depict your implementation, if necessary.
- Delete italic text when submitting your report. Those are only guidelines to help you
- The name of the file must match *201X-XXXXX_kernellab_report.pdf* and placed under *report* directory.

6. Submission Guidelines

Submission

- **Due Date**
 - Tue., October 15, 16:59
- **Submission Files**
 - Source code
 - Makefile
 - Report (in PDF format, file name: 201X-XXXXXX_kernellab_report.pdf)
- Follow the report template and naming convention
- Do not change the directory structure
- Do not email us your code and report
- You can add more files (.c, .h) for your implementations but the module should be built by using only the *make* command
- Failure to follow any of the submission guidelines can result in a **deduction** of your score

Note

- This lab is quite **tough**
 - You need to look through the kernel source
 - Debugging for the kernel is not an easy task
- **So, start working as soon as possible**

References

- **Linux Kernel Module Programming Guide**
 - <http://www.tldp.org/LDP/lkmpg/2.6/html/>
- **Debugfs APIs**
 - <https://www.kernel.org/doc/Documentation/filesystems/debugfs.txt>
- **Address Translation & Page Table Entry**
 - Class Text Book: “Computer Systems A Programmer’s Perspective, Randal E. Bryant, David R. O’Hallaron, 3rd International Edition, Pearson, 2016” - page 862
 - Class Text Book: “Computer Systems A Programmer’s Perspective, Randal E. Bryant, David R. O’Hallaron, 3rd International Edition, Pearson, 2016” - page 863 ~ 864
 - Intel® 64 and IA-32 Architectures Software Developer’s Manual, Volume 3, Chapter 4
- **Makefile Guide**
 - <https://www.cs.duke.edu/~ola/courses/programming/Makefiles/Makefiles.html>